Strengths

- ✓ Only one layer changed, 512×20×20 interface kept
- ✓ Selective global context: 4-direction scan
- ✓ Dedicated shape path (1×5, 5×1 kernels and gate)
- ✓ Stable start: γ = 0.001, inherits cv1, cv2, FFN
- ✓ PyTorch only: CPU/GPU, deterministic

Limitations

- ✗ Parameters +7.17%; larger FP32 checkpoint
- ✗ CPU latency about 2.1×
- ✗ Training about 2.3 → 4.7 h per run
- ✗ O(L log L) scan, many intermediate tensors
- ✗ At P5 (L = 400), the SSM advantage does not show